


Source discovery → measurable learning loop

9 tickets · 3 workstreams

Status: planned, not built

Truth: GitHub #17–#25

Research tool: localhost only



Arrow = blocker. Parallel starts: 01, 04, 05, 06, 07.

Goal & audience

- **Goal:** reach first users fast, get feedback, turn creative examples into testable content/ad hypotheses. No audience yet.
- **Audience (confirmed):** students/new grads prepping SWE interviews; can attempt Python; want to practice explaining, testing, revising.
- **Positioning (hypothesis):** “Practice explaining your Python solution, see what to improve, try again.”
- **Claims:** Python voice practice (Mock/Coach), evidence-linked review, retries. Samples labelled, measured separately.

Tickets

#	TICKET	OUTCOME	NEEDS
01	Discover + classify · #17	Search → Exa/Firecrawl/Monid → Gemini → Jev	—
02	Related creatives · #18	Embeddings retrieve sources + evidence	01
03	Strategy + scripts · #19	Evidence → original hooks, scripts, reviewed briefs	02
04	Landing + waitlist · #20	Explain offer; truthful next action	—
05	Contextual feedback · #21	Preset answers + optional text; triage	—
06	Bug reporting · #22	Sitewide reports, repro, resolution flow	—
07	Analytics + heatmaps · #23	Events, funnels, safe heatmaps	—
08	Landing experiments · #24	Variants vs. approved hypothesis + baseline	04, 07
09	Content planning · #25	Outcomes + feedback → briefs	03, 05, 07

Content strategy

PROBLEM	CONTENT	PROOF	ACTION
Can't explain reasoning	Explain while coding	Real/labelled demo	Try practice
Doesn't know what to fix	Specific finding on attempt	Evidence citation	Attempt + read feedback
Misses cases / repeats	Fail → focused retry	Before/after	Practice, rate usefulness

- Clips = one moment; long-form = full loop. Never copy others' scripts.
- Signals: visits, attempt start/finish, review opened, retry, user comments. Popularity/model scores ≠ conversion.

WS1 · Dev research (01–03)

Localhost only. Loopback-bound. No prod routes/assets/jobs. Keys + data outside deploy. Hidden URL ≠ isolation.

search → Exa·Firecrawl·Monid → media → Gemini evidence → Jev labels → embeddings → LLM analysis → human brief → outcome

- **Exa** discovery · **Firecrawl** pages · **Monid** Meta/TikTok ad collectors
- **Gemini** observations · **embeddings** retrieval · **Jev** classify · **LLM** (any) writes scripts
- User starts runs; topic or blank (uses saved context). No scheduled collection.
- Unknown = unknown; never fabricate. Uncertain → review.
- Records: Source, Evidence, Judgment, Brief — versioned, cited.

- Failures resumable; cost/latency logged.
- No customer data in research; only approved aggregates.
- Human review before publish. No auto ad spend/posting.

WS2 · Site signals (04–07)

- **Landing:** audience/CTA → walkthrough → benefit → availability → privacy/FAQ → waitlist. Email + purpose, dedupe, withdraw. No fake proof/urgency.
- **Feedback:** collapsed panel, preset + optional text, never interrupts task.
- **Bugs:** expected/actual, repro, allowlisted metadata. No audio/code/transcripts/secrets. No auto GitHub posts.
- **Analytics:** event allowlist. Never record interview content; unsafe surfaces excluded from heatmaps.
- Each: own records, operator view, coverage matrix. Analytics down ≠ feedback/bugs down.

WS3 · Experiments (08–09)

- **Tests:** stable assignment, predeclared outcome, guardrails, bot exclusion, rollback. Split/size/duration from baseline, not gut.
- **Planning:** aggregate outcomes + feedback → briefs. Personal data separate.

Open owner decisions

- **Launch CTA** — waitlist vs. practice (#8)
- **Distribution** — channels, organic vs. paid (#9)
- **Data rules** — consent, retention, contact; privacy text says no analytics
- **Paid-run scope** — budgets, limits
- **Experiment method** — stats + thresholds
- **Providers** — email, heatmaps

Out of scope

Hosted research tool, scheduled/bulk scraping, publishing, ad spend, fake proof, new languages, guaranteed results.